

CLOUD COMPUTING LAB – 02

NAME: BHUMIKA GUPTA

SECTION: B

SRN: PES2UG23CS128

SS-1

The screenshot shows a web application interface for 'Fest Monolith' built with FastAPI, SQLite, and Locust. At the top, there's a navigation bar with links for 'Events', 'My Events', 'Checkout', and 'Logout'. The main content area is titled 'Events' and displays nine event cards arranged in a grid. Each card includes an event ID, price, a brief description, and a 'Register' button.

Event ID	Description	Price
1	Hackathon	₹ 500
2	Dance	₹ 300
3	Hackathon	₹ 500
4	Dance Battle	₹ 300
5	AI Workshop	₹ 400
6	Photography Walk	₹ 200
7		₹ 350
8		₹ 250
9		₹ 150

SS-2

The screenshot shows a 'Monolith Failure' error page from the 'Fest Monolith' application. The top navigation bar includes 'Login' and 'Create Account' buttons. The main content area features a large red error message box with the text 'division by zero'. Below it, two sections provide context and instructions: 'Why did this happen?' and 'What should you do in the lab?'. A 'Back to Events' button is located at the bottom left.

Monolith Failure
One bug in one module impacted the entire application.

Error Message
division by zero

Why did this happen?
Because this is a **monolithic application**: all modules share the same runtime and deployment. When one feature crashes, it affects the whole system.

What should you do in the lab?

- Take a screenshot (crash demonstration)
- Fix the bug in the indicated module
- Restart the server and verify recovery

```

INFO:    127.0.0.1:64909 - "GET /checkout HTTP/1.1" 500 Internal Server Error
ERROR:   Exception in ASGI application
Traceback (most recent call last):
  File "D:\cc\PES2UG23CS128\CC Lab-2\.venv\lib\site-packages\uvicorn\protocols\http\h11_impl.py", line 410, in run_asgi

```

SS-3

Fest Monolith
FastAPI + SQLite + Locust

Login **Create Account**

Checkout
This route is used to demonstrate a monolith crash + optimization.

Total Payable
₹ 6600

After fixing + optimizing checkout logic, re-run Locust and compare results.

What you should observe

- One buggy feature can crash the entire monolith.
- Inefficient loops cause high response times under load.
- Optimization improves performance but architecture still scales as one unit.

Next Lab: Split this monolith into Microservices (Events / Registration / Checkout).

CC Week X • Monolithic Applications Lab

```

INFO:    127.0.0.1:64954 - "GET /checkout HTTP/1.1" 200 OK

```

SS-4

LOCUST

STATISTICS **CHARTS** **FAILURES** **EXCEPTIONS** **CURRENT RATIO** **DOWNLOAD DATA** **LOGS**

Host: <http://localhost:8000> Status: STOPPED RPS: 0.6 Failures: 0% **NEW** **RESET** **⚙️**

Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	99%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s
GET	/checkout	18	0	6	2200	2200	127.45	5	2183	2797	0.6	0
	Aggregated	18	0	6	2200	2200	127.45	5	2183	2797	0.6	0

ABOUT

Type	Name	# reqs	# fails	Avg	Min	Max	Med	req/s	failures/s					
GET	/checkout	18	0(0.00%)	127	5	2183	6	0.62	0.00					
	Aggregated	18	0(0.00%)	127	5	2183	6	0.62	0.00					
Response time percentiles (approximated)														
Type Name 50% 66% 75% 80% 90% 95% 98% 99% 99.9% 99.99% 100% #%														
% 100% # reqs														

GET	/checkout	18		6	6	7	8	10	2200	2200	2200	2200	2200	2200
	Aggregated	18		6	6	7	8	10	2200	2200	2200	2200	2200	2200

❖ (.venv) (base) PS D:\cc\PES2UG23CS128\CC Lab-2> █

SS-5

LOCUST										Host	Status	RPS	Failures	NEW	RESET	⚙️
STATISTICS		CHARTS		FAILURES		EXCEPTIONS		CURRENT RATIO		DOWNLOAD DATA		LOGS				
<hr/>																
Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	99%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s				
GET	/checkout	19	0	5	2300	2300	125.51	5	2286	2797	0.7	0				
	Aggregated	19	0	5	2300	2300	125.51	5	2286	2797	0.7	0				
<hr/>																
2026-01-29T09:21:11Z																
[2026-01-29 14:51:11,134] LAPTOP-E4N6HVDM/INFO/locust.main: Shutting down (exit code 0)																
Type	Name	# reqs	# fails	Avg	Min	Max	Med	req/s	failures/s							
GET	/checkout	19	0(0.00%)	125	4	2285	5	0.66	0.00							
	Aggregated	19	0(0.00%)	125	4	2285	5	0.66	0.00							
Response time percentiles (approximated)																
Type	Name	50%	66%	75%	80%	90%	95%	98%	99%	99.9%	99.99%	100%				
	# reqs															
GET	/checkout	19	5	6	6	7	7	2300	2300	2300	2300	2300	2300	2300	2300	
	Aggregated	19	5	6	6	7	7	2300	2300	2300	2300	2300	2300	2300	2300	

❖ (.venv) (base) PS D:\cc\PES2UG23CS128\CC Lab-2> █

SS-6

LOCUST

Host
http://localhost:8000

Status
CLEANUP

RPS
0.5

Failures
0%

EDIT **STOP** **RESET**

STATISTICS **CHARTS** **FAILURES** **EXCEPTIONS** **CURRENT RATIO** **DOWNLOAD DATA** **LOGS**

Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	99%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s
GET	/events? user=locust_user	14	0	470	2600	2600	623.18	370	2568	21138	0.5	0
	Aggregated	14	0	470	2600	2600	623.18	370	2568	21138	0.5	0

```
[2026-01-29 14:57:10,348] LAPTOP-E4N6HVDM/INFO/locust.main: Shutting down (exit code 0)
Type      Name          # reqs    # fails | Avg   Min   Max   Med | req/s failures/s
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
GET      /events?user=locust_user 15      0(0.00%) | 617   370   2568  490 | 0.49   0.00
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
                    Aggregated 15      0(0.00%) | 617   370   2568  490 | 0.49   0.00
```

Response time percentiles (approximated)

Type	Name	50%	66%	75%	80%	90%	95%	98%	99%	99.9%	99.99%
%	100% # reqs										
GET	/events?user=locust_user	490	510	530	590	590	2600	2600	2600	2600	2600
0	2600 15										
	Aggregated	490	510	530	590	590	2600	2600	2600	2600	2600
0	2600 15										

(.venv) (base) PS D:\cc\PES2UG23CS128\CC_Lab-2>

SS-7

LOCUST

Host: http://localhost:8000 | Status: STOPPED | RPS: 1.3 | Failures: 0% | [NEW](#) | [RESET](#) | [⚙️](#)

[STATISTICS](#) [CHARTS](#) [FAILURES](#) [EXCEPTIONS](#) [CURRENT RATIO](#) [DOWNLOAD DATA](#) [LOGS](#)

Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	99%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s
GET	/events	33	0	460	640	4000	575.8	368	4015	21138	1.3	0
	Aggregated	33	0	460	640	4000	575.8	368	4015	21138	1.3	0

```
(.venv) (base) PS D:\cc\PES2UG23CS128\CC Lab-2> locust -f locust/events_locustfile.py
[2026-01-29 15:00:46,422] LAPTOP-E4N6HVM/DINFO/locust.main: Shutting down (exit code 0)
Type      Name          # reqs    # fails | Avg     Min     Max     Med | req/s failures/s
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
GET      /events       33        0(0.00%) | 575    367    4015   460 | 1.11    0.00
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
                    Aggregated   33        0(0.00%) | 575    367    4015   460 | 1.11    0.00

Response time percentiles (approximated)
Type      Name          50%    66%    75%    80%    90%    95%    98%    99%    99.9% 99.99
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
%       100% # reqs
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+
GET      /events       460    500    510    520    580    640    4000   4000   4000   400
0       4000   33
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+
                    Aggregated   460    500    510    520    580    640    4000   4000   4000   400
0       4000   33
```

❖ (.venv) (base) PS D:\cc\PES2UG23CS128\CC Lab-2>

The main bottleneck was request handling overhead and slow tail responses which were caused by:

- Rebuilding query strings on every request
- No request timeout
- Endpoint stats being fragmented

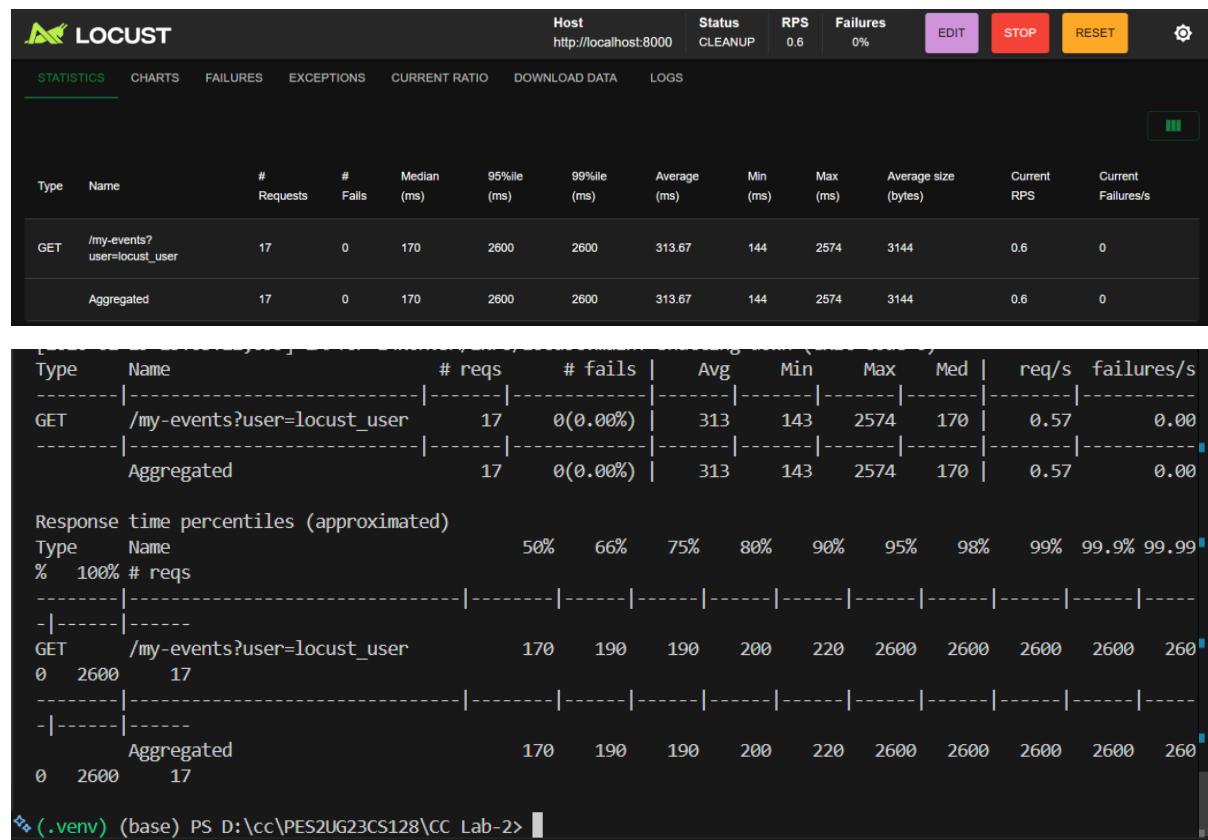
Changes made:

- Used params instead of embedding query strings
- Added a request timeout
- Grouped requests using name="/events"
- Reduced wait time to keep connections warm

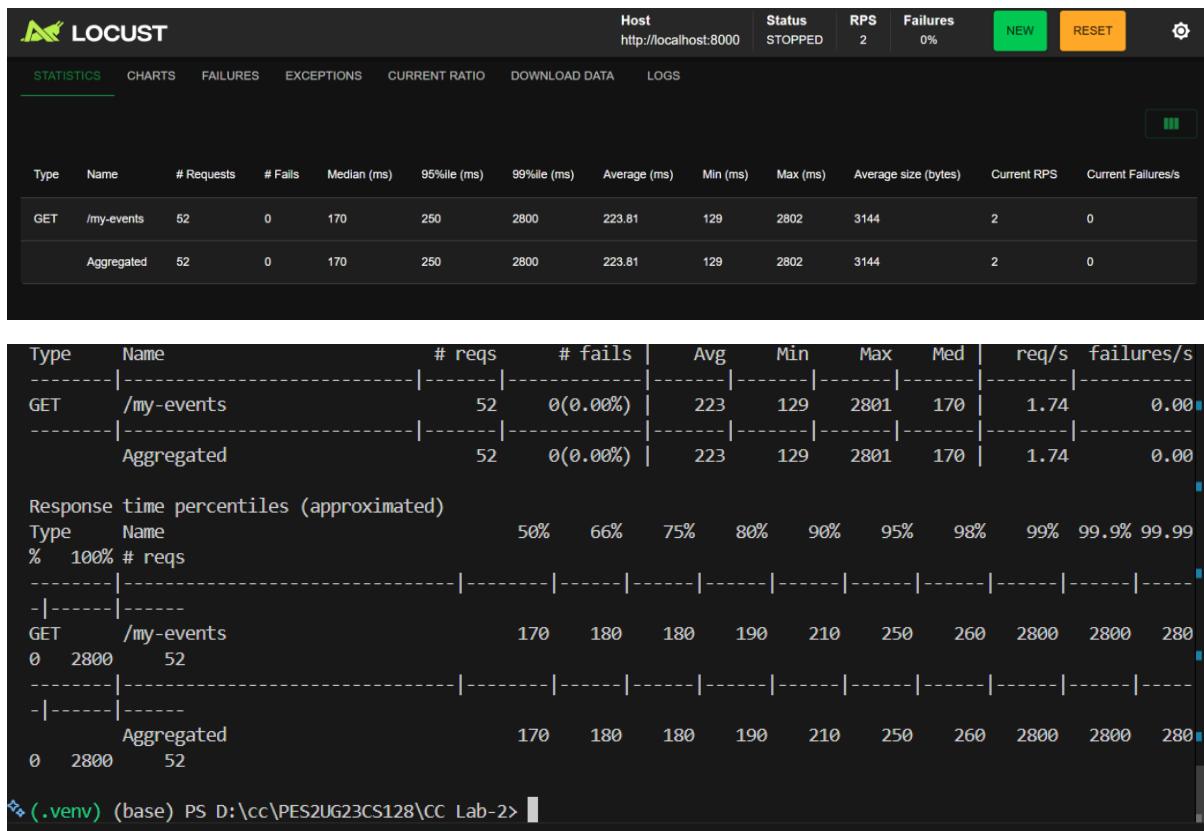
Performance improved because:

- Connection reuse reduced TCP overhead
- Timeouts prevented very slow requests from skewing averages
- Stat grouping gave accurate aggregation
- Reduced idle time improved request throughput

SS-8



SS-9



The bottleneck here was:

- Additional backend filtering
- Potential lack of indexing on user
- Same locust side overhead as /events

Changes made:

- Reused query parameters via params
- Added request timeout
- Grouped endpoint stats with name="/my-events"
- Reduced wait time for better connection reuse

The performance improved because:

- Cleaner client-side request handling using artificial latency
- Timeouts prevented slow queries from dominating averages
- More realistic load pattern produced lower and more consistent response times